

# Pneumonia Classification

Pneumonia classification from chest X-rays using deep learning

Hakim Chekirou

---




# Why does it matter ?


- Pneumonia accounts for over 15% of all deaths of children under 5 years old.
- It requires highly trained specialists and confirmation through clinical history.
- Other conditions such as fluid overload (pulmonary edema), bleeding, volume loss, lung cancer, or post-radiation or surgical changes make the prediction harder.
- A number of factors such as positioning of the patient and depth of inspiration can alter the appearance of the CXR.
- In addition, clinicians are faced with reading high volumes of images every shift.

# Dataset

- Pneumonia detection challenge
- 26,684 Chest X rays in the form of Dycom files.
- Unbalanced dataset (20672 negative vs 6012 positive examples).
- High resolution 1024\*1024 in gray scale



Positive example



Negative example

# Data preparation


## Preprocessing

- Normalizing.
- Train and test split (90 %-10%).

## Data Augmentation :

- Random rotation up to  $6^\circ$
- Random horizontal flipping

# Convolution networks (alexnet)




- Pretrained on imagenet
- Finetuned for 20 epochs
- SGD and weighted Cross-entropy loss

Alexnet architecture

Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.89	0.88	0.89	76.64%	83%
Positive	0.62	0.65	0.64		


Alexnet results

# Residual networks



Residual block

- Pretrained on ImageNet and finetuned for 20 epochs.
- Using weighted cross entropy loss
- and SGD with momentum



Resnet50 architecture

# Residual networks results


Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.90	0.87	0.88	76.84%	82 %
Positive	0.60	0.67	0.63		

Results for the Resnet50 architecture

Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.94	0.75	0.84	79.71%	77 %
Positive	0.50	0.84	0.63		

Results for the Resnet18 architecture

# Inception v3



- Pretrained on imagenet
- Finetuned for 20 epochs
- SGD and weighted CE loss

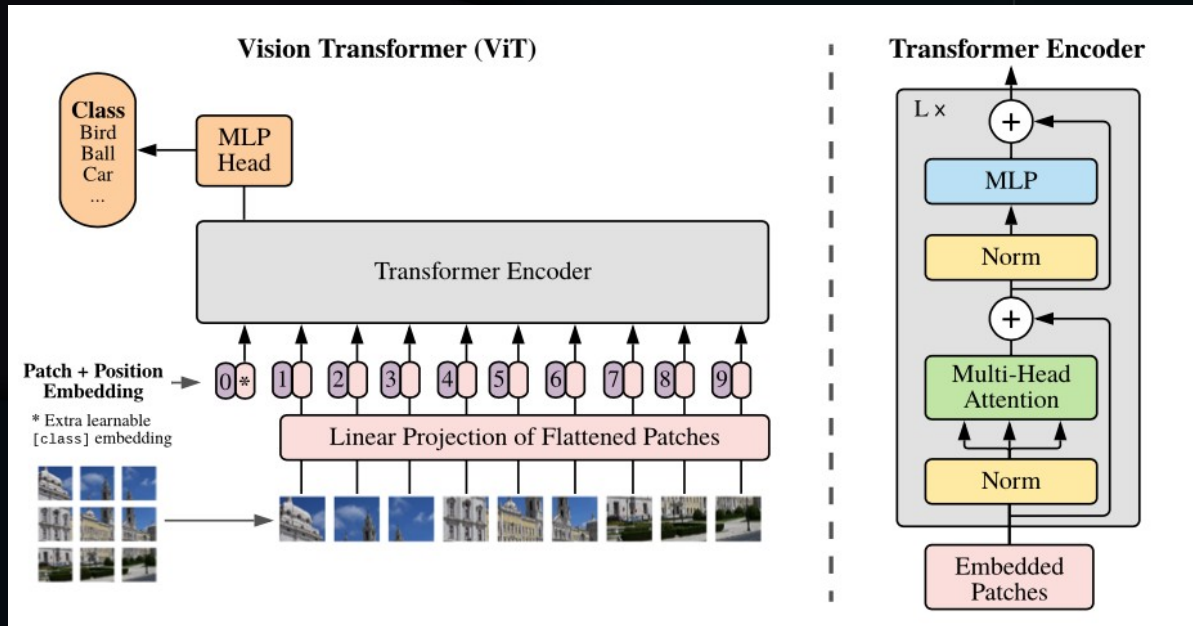
Inception v3 architecture

Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.91	0.84	0.88	78.64%	82%
Positive	0.58	0.73	0.65		

Inception v3 results



# Proposition : Vision Transformers




- Trained from scratch using two sets of parameters.
- Variation of number of patches : 64 vs 256
- Use of a CLS token or mean representation for classification.
- Depth fixed to 6 layers
- MLP dimension of 1024
- Use of dropout in the attention weights and positional embeddings.

An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale, Dosovitskiy et al, 2020.

# Vision Transformer results

Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.89	0.75	0.82	72.51%	74%
Positive	0.46	0.70	0.55		

Interpretability : visualizing the attention weights



# Ensemble Model

Combining Predictions from all the mentioned models :  
mean prediction of the 5 models considered.

Class	Precision	Recall	F-1 score	ROC AUC	Accuracy
Negative	0.92	0.85	0.88	79.28%	82 %
Positive	0.59	0.74	0.66		

Similar results to Resnet18 in terms of ROC AUC score but with better accuracy.

# Conclusion

---

- The best ROC AUC score is achieved with the resnet18 model (79.71%).
- The ensemble model achieves a similar score in terms of ROC AUC (79.28%) Score but with a better accuracy and is thus preferred.
- The transformer model is more interpretable.